Cardiac Arrhythmia Disease Classification Using LSTM Deep Learning Approach

Muhammad Ashfaq Khan and Yangwoo Kim

Computers, Materials & Continua 2021, 67(1), 427-443.

Paper Review by Taha Cheema - 21204729

This paper proposes a deep learning pipeline for classifying cardiac arrhythmias from electrocardiogram (ECG) records into one of 16 disease categories. To build their classifier, the authors adopt a three-stage pipeline on a publicly available arrhythmia dataset (from the UCI Machine Learning Repository, a widely used hosting site for benchmark datasets), which contains 452 patients each described by 279 medical features. First, they preprocess the data by normalizing feature values and dropping rows and columns with missing or zero entries. Second, they apply principal component analysis (PCA) to reduce the number of features, cutting noise and shrinking the input space. Third, they feed the reduced features into a Long Short-Term Memory (LSTM) network followed by a few dense layers and a softmax output. To handle the heavy class imbalance in the dataset—where most records are normal ECGs—they randomly duplicate minority-class examples before training. The authors' concluding remarks are that the PCA + LSTM hybrid achieves higher classification accuracy than prior machine learning approaches on the same dataset and is well-suited to high-dimensional medical data.

A major critique of this paper is that its abstract claimed an accuracy of 93.5% on 16-class classification but in the actual experiment, they only used 6 classes. This detail is discretely placed in Section 4 where the authors write that "classes which contained an irrelevant number of instances were eliminated". It is unclear what they mean by irrelevant number of instances. Secondly, the authors' approach to handle the class imbalance is not ideal for two reasons: 1) the model could overfit to a very specific set of examples in rarer classes; 2) the paper does not state whether duplicates were generated before or after the test/train split of the data. If duplicates were generated first and then the dataset was split, the same patient record can land in both training and test sets, inflating reported accuracy.

However, beyond these specific critiques, it is important to place this paper in the broader landscape of AI and medicine. Healthcare is perhaps the most important scientific pursuit for humanity and with the advent of AI, there is immense potential for improving the field. Classifications tasks in medicine, such as the one discussed in this paper, are excellent potential application areas for AI because it can speed up initial screenings and perhaps even help triage cases better in a medical system that is often facing a shortage of health practitioners in many countries. Disease screening and classification of early markers of certain diseases through AI analysis of ECGs or general radiology artifacts can also help democratize access to healthcare as as health specialists access may be monetarily infeasible for certain strata of the population in some regions. But with these opportunities, we need to be extremely careful and avoid pitfalls like biases, that would at best erode trust for clinical use of AI and at the worst, lead to wide-scale incorrect diagnoses. This paper's authors designed a solution to solve a life-critical clinical problem (arrythmias) but their methodological choice of dropping rare (or "irrelevant") classes is precisely what could lead to false negative classification of under-represented arrythmias if this system was deployed in the real-world. Therefore, this is perhaps the biggest takeaway from this paper: great potential for benefits in soceity but grave dangers from incorrect training and application.

Lastly, I would like to reference a remark I added in my review of Matt Welsh's paper, 'The End of Programming', where I mentioned that programming cannot be abstracted away completely in a world of AI coding agents because safety-critical software cannot withstand potential errors from LLMs. Similarly, taken at their best, this paper's best recall value of 90.7% is not sufficient for standalone medical use because false negative diagnoses of even a few percent can have big consequences. This does not mean that any medical AI research not having perfect recall scores is meaningless, it just means that such research must be complemented with more HCI focused research about how humans can best *collaborate* with AI in medical settings instead of blindly trusting. Such research could help improve diagnostic accuracy of health practitioners while minimizing the chances of being misled as collateral. For example, a recent trial showed how color-coding LLM diagnostic advice for doctors based on the degree of agreement between different LLMs on the same prompt can be the right kind of nudge for practitioners to trust AI at the right moments and pause and scrutinize at others (Qazi et al.).

References

Ihsan Ayyub Qazi, Ayesha Ali, Asad Ullah Khawaja, Muhammad Junaid Akhtar, Ali Zafar Sheikh, Muhammad Hamad Alizai, Mitigating Automation Bias in Physician-LLM Diagnostic Reasoning Using Behavioral Nudges: A Randomized Controlled Trial, medRxiv 2026.06.01.26354596; doi: https://doi.org/10.64898/2026.06.01.26354596